

Process of Compilation

Program development proceeds as follows:

1. Edit the program in some text editor **vi,vim,gedit,emacs**.
The result is a text file.
2. Compile this text file into machine code.
3. Run the resulting machine code.

Interpretation versus Compilation

Some languages, most notably Java and *C#* use **interpreters**.

The compiler does not compile into native machine code, but into an intermediate code. (Java Byte Code).

1. Programmer writes program in some text editor.
2. Compiler translates into Java Byte Code.
3. Interpreter (Java Virtual Machine) reads the Java Byte Code, and executes its instructions. (You can guess that the JVM consists mostly of 'ifs' of the form 'if next instruction is ... then do ...'. Interpretation is necessarily slower than native code. It also uses more memory, and more energy.
4. Modern implementations of JVM combine interpretation with compilation. (Just in time compilation)
5. Main advantage of interpretation is portability.

Separate Compilation

- Compilation can be quite time consuming for big programs, even on a fast computer.
- Different parts of a program are written by different people.
- Some contributors may want to hide their sources from you.

Because of this, we have **separate compilation**.

Separate Compilation (2)

- Program consists of different text files, which can be edited separately.
- For each of the files separately, an **object file** is created. (They have extension .o in Linux, .obj in Windows.) You can use `nm filename.o --demangle` to see what is inside an object file.
- The linker collects all object files, fills in the cross references (mostly addresses of function calls) and creates a single executable file. (No extension in linux, .exe extension in Windows.)

Separate Compilation (3)

Unfortunately, completely separate compilation is impossible.

File **rational.cpp** defines rational numbers and operations on them.

Another file, **matrix.cpp** defines a matrix as an array of rationals, and defines operations on matrices.

Main file uses matrices and rationals.

- When producing a local variable of some user defined type, the compiler needs to know how much space to reserve.
- When compiling `v.x`, the compiler needs to know that `v` has a field `x`, and which type it has. (and the offset if also useful)
- When compiling a function call `f(m1, ..., mn)` the compiler needs to know that `f` exists, and what types it expects and returns.

file.h versus file.cpp

Every program file must be split into two files: **file.h** contains:

- Declarations of available types (classes and structs).
- Declarations of available functions.
- Definitions of small functions that are **inlined**. Inlining means that the compiler replaces the function call by its definition. This is useful for small functions that are called very often.

The rest stays in file **file.cpp**.

Including

The method for reading the necessary `.h` files is as primitive as you can imagine. Write `#include "file.h"` at the top of your `.cpp` file, for every file whose declarations you need. (For system defined files, it is `#include <library>`)

If you have nested dependencies, then it is quite possible that the same file gets included twice, which is not good. (The number can easily grow exponentially.)

(For example, `main` may use `rational.h` and `matrix.h`, which in turn uses `rational.h`.)

In order to avoid this, `include guards` are used.

```
#ifndef MATRIX_INCLUDED
#define MATRIX_INCLUDED 1
    // Name must be unique and in capitals.
#include "rational.h"
struct matrix
{
    rational repr[2][2];
};

matrix operator * ( matrix m1, matrix m2 );
matrix operator * ( rational r, matrix m );
matrix operator + ( matrix m1, matrix m2 );
    // Declare a lot of operators.
};
#endif
```


Linking

The linker collects all the .o files, and makes a complete program from it.

It looks for functions that are called in one of the files, and tries to find a definition in one of the other files.

If the linker cannot find a definition, you will get an error message that is rather unpleasant:

```
test.o: In function 'rational::rational(int, int)':
test.cpp:(.text._ZN8rationalC2Eii[_ZN8rationalC5Eii]+0x2d):
      undefined reference to 'rational::normalize()'
collect2: error: ld returned 1 exit status
```

Possible Causes of Link Errors

Link errors appear when you declare a function in one of the **.h** files (the compiler now believes that the function exists), but don't define it in any of the **.cpp** files.

It may also happen that you forgot to pass one of the **.o** files to the linker.

Another possibility is that the type of the definition slightly differs from the type of the declaration. (Usually wrong **constness**, or different **namespace**).

A third possibility is when a function is defined multiple times. This usually happens when you forget to write **inline** in a function defined in a **.h** file.

Make

If you edit a .cpp file, then only its corresponding .o file needs to be recompiled.

If you edit a .h file, only the the .o files of the .cpp files that `#include` it, need to be recompiled.

`make` sorts this out automatically.

Makefile

```
program:  f1.o f2.o f3.o
          g++ -o program f1.o f2.o f3.o

f1.o: f1.cpp f2.h f3.h
      g++ -c f1.cpp -o f1.o -std=c++14

f2.o: f2.cpp f1.h f3.h
      g++ -c f2.cpp -o f2.o -std=c++14

... (more for the other obj files.)
```

Specify for each file that the compiler/linker constructs, the files from which it is constructed.

After that, give the command that does the construction. (The 8 spaces are a Tab.)

The linker is also called g++.

Revision/Version Control

In case more than one person works on the same project, different people may have different versions of the sources.

Keeping the sources consistent may get pretty difficult.

Usually, one programmer does not want that others use unfinished versions of his sources. So he makes a local copy, works on some of the sources, and shares his sources only when they are finished.

This is called **commit**.

Since many programmers may be doing this at the same time, it is useful to automate the revision control.

Revision Control Systems

- keep previous versions of the sources, for the case that some improvement turns out not an improvement, or an error is introduced.
- keep different versions of sources, e.g. one for experimenting, and one reliable versions. Sometimes the complete program is sold in different versions.
- check when different users commit conflicting changes. Some systems automatically or semi-automatically merge the differences.

I do not have any experience with such systems, but you must know that they exist. Often used systems are **Git** and **CVS**.

Types of Variables

C^{++} has many types of variables, which is often confusing. We have seen before that the default behaviour for assignment, initialization and parameter passing in C^{++} is **copying**. This ensures that different variables are independent of each other, which makes the program easier to understand, or to analyze automatically. (This is **value semantics**).

Reference Variables

Sometimes you do not want to copy for one of two possible reasons:

(1) A procedure must be able to change a variable, and (2) copying large objects is inefficient.

A **reference** is a short-lived variable that does not have contents of its own, but which shares its contents with another variable.

In case 1, use **reference**, in case 2, use **const reference** or **rvalue reference**.

There is no equal sharing. One (value) variable is the owner, and the references exist shorter than this variable. Therefore, no garbage collection is needed.

Use of References in Parameter Passing

```
matrix operator + ( const matrix& m1, const matrix& m2 );  
    // There is no need to make a copy of matrix m1, m2  
    // in order to add them. The 'const' keyword indicates  
    // that the reference will never change the value  
    // it refers to.
```

```
std::ostream& operator << ( std::ostream& stream,  
                             const matrix& m );  
    // Similarly, there is no need to copy a matrix  
    // in order to print it.
```

If you have a function parameter that could be copied in principle, but you worry about efficiency, use `const&`.

Use of References in Parameter Passing

```
void operator += ( matrix& m1, const matrix& m2 );  
    // The += operator adds the second matrix to the  
    // first. The second parameter could be just 'm2'  
    // but that would be inefficient, so we made it  
    // const reference.  
    // The first parameter must be a reference,  
    // because += must be able to change it.
```

If you want a function to be able to change something through one of its parameters, then use `&` without `const`.

Note that nearly always it is better to return a value, than to modify a parameter. If you decide to modify a parameter, it must be very visible from the function name. (Above it is `+=`).

Rvalue References

Dilemma:

```
matrix inverse( const matrix& m );  
    // No copying. Efficient, but we cannot change m  
    // during computation.  
matrix inverse( matrix m );  
    // Copying. More costly, but we can use m as  
    // scratch area during computation.
```

Another type of reference: We take a reference to some variable, and the reference is the last user of the current value of the variable.

```
matrix inverse( matrix&& m );  
    // Not copied, but we can use m as scratch.  
    // Nobody cares, because we are its last user.
```

Rvalue References (2)

Rvalue references can be used when main variable either gets overwritten, or goes out of scope:

```
{  
    matrix m1 = ...  
    std::cout << inverse( m1 );  
    // m1 is going to be overwritten.  
    {  
        matrix m2 = ...  
        m1 = inverse( m2 );  
        // m2 goes out of scope.  
    }  
}
```

Don't worry about Rvalue references now. We will come back to them later.

C

Note that *C* also has references, but they are invisible:

```
a = b;           // Possible.
a.x = b.x        // Possible.
a[2] = b[3];     // Possible.
*x = y;          // Possible.

a+4 = x;         // Not possible.
f(1) = x;        // Not possible.
&(a+2) = x;      // Not possible.
```

In *C*, every expression has a type and in addition, is either an **lvalue** or **rvalue**.

The name derives from the question: Can the expression occur left of an assignment?

Lvalues represent addresses (are a kind of implicit pointers).

Rvalues represent values.

Variables, operator `.` (record selection), operator `[]` (array indexing), and operator `*` (pointer dereference) create lvalues.

Other expressions are rvalues.

If a function requires an **rvalue**, but the argument is an **lvalue**, the compiler inserts a **load**:

```
a = b.x + y[i] + 5;
```

```
a = ( load( b.x ) + load( y[ load(i) ] ) ) + 5;
```

C-style lvalues are *C*⁺⁺-references.

C^{++} made the references accessible to the user, and added more types of references:

	creator of reference	receiver of reference
<code>X&</code>	wants to receive information	wants to send information
<code>const X&</code>	wants to send information	wants to receive information, won't change information during computation
<code>X&&</code>	wants to send information, does not need information anymore, after sending	wants to receive information, may change its value during computation.

Type `X&&` is called **rvalue reference**. (Because it replaces `const X&`, which in turn replaces a copy.)

If you need

creator: wants to send information, and still use same information later.

receiver: wants to receive information, may change its value during computation,

you will need to make a copy.

As for non **const X&**, use it rarely. Instead, try to:

1. return a value through a function:

```
something = value( val );
```

2. modify through member function.

```
something. set_value( val );
```

Both are better than `set_value(something, val);`

Iterators

An **iterator** shares with an element of a container containing more than one element. (array, hashmap, vector, list). Iterators are a convenient way of accessing the elements of a container. For many container types, they are the preferred way of reaching the elements.

```
std::vector<int> vect = { 1,2,3,4,5 };

for( std::vector<int> :: iterator p = vect. begin( )
    p != vect. end( );
    ++ p )
{
    *p = *p + 1;
}
```

Iterators (2)

An `std::list` is implemented completely different than an `std::vector`. The user doesn't see that:

```
std::list<int> vect = { 1,2,3,4,5 };

for( std::list<int> :: iterator p = vect. begin( );
    p != vect. end( );
    ++ p )
{
    *p = *p + 1;
}
```

Pointers

A **pointer** is the most general form of sharing variable (and also the most dangerous).

- It can be **nullptr**. (not refer to anything at all)
- It can be used for allocating space on the heap, i.e. **own** the object (not share with a local variable).
- A pointer can be changed during its life time, to start pointing somewhere else, between different containers or variables.

Pointers are essential for building iterators and for defining complicated objects.

Don't use pointers for anything else.

In the latter case, the object must restore value semantics through defined **copy-constructor** and **assignment operator**. Always aim at value semantics for defined types.

struct versus class

Object-Oriented Programming is popular, but nobody can define what it is.

I think the main point that in OOP there are **classes**, while C has only **struct**.

Both struct and class are tools for definition of user types.

A class can hide implementation which makes it easy to maintain complicated invariants and equivalences for the user of a type.

With struct, the user has to remember the invariants and equivalences of the type.

Classes

Class = Representation + Invariants + Equivalences.

Suppose that we want to implement rational numbers of form $\frac{p}{q}$.

A rational number can be represented by a pair of integers (p, q) .

(p, q) and (pn, qn) represent the same number.

- Either we use an invariant, that p, q have no common factors and $q \geq 0$. In that case we do not need equivalences.
- Or we use no invariant, and we design our class in such a way that (p, q) and (pn, qn) are indistinguishable. In that case we do not need an invariant.

Invariant: Not every possible representation represents an object of the new type.

Equivalence: Different representations represent the same object in the new type.

Dates: Suppose we want to implement dates:

A data can be implemented as a triple

$\text{Nat} \times \{1, \dots, 12\} \times \{1, \dots, 31\}$.

Not every combination can occur, some of the rules are quite complicated.

In this case, the invariant is the combination of rules that state which combinations are a valid date.

In C , the **struct** is always open, and the user of the **struct** needs self control, and to remember that there is an invariant.

In C^{++} , access can be restricted to a small set of methods that preserve the invariant. The user of the **class** doesn't need to remember anything.

I think that this is the main advantage of **class** over **struct**. I will not try to define what OOP is.

Constructors, Controlled Access

Modern languages (like C^{++} , Java, Python) help the programmer by means of constructors and controlled access.

Constructors make sure that you cannot construct representations that break the invariant. Controlled access make sure that you cannot distinguish representations that are supposed to be equal.

Of course, a *real programmer* can keep track of all these invariants and equivalences in his head.

But doing this for 10 types at the same time for classes written by somebody else, or by even yourself one year ago, will be a challenge, even to the best programmer.

Constructors

Methods with the same name as the class they occur in, are called **constructors**. The task of the constructors is to establish the class invariants.

If you provide constructors for a class, it is impossible to obtain class objects in other ways. Every function that constructs a class object, will have to do this by means of a constructor.

If you write a few constructors, and manage to get them right, class invariants for new objects are guaranteed.

Constructors can have any types of arguments, but a few constructors are special:

- The **default constructor** has no arguments. It is automatically inserted when a variable is declared without initializer.
- A **copy constructor** has one argument of type `const C&`, `C&&`, or `C&`. It is automatically inserted when a parameter is transferred by value to a function.

The compiler may also decide to use it, when a function returns a variable.

Controlled Access

- Allow only a small set of functions to see the representation. If you manage to ensure that those functions do not distinguish equivalent objects, then no other function can.
- Allow only a small set of functions to modify the representation (together with the constructors). If you manage to ensure that these functions preserve equivalence and the invariants, then all functions must do this.

Destructors

Destructors are special. They do preserve invariants, but not the class invariants. They preserve global invariants, usually resource allocation invariants:

‘Everything that is allocated is connected to a local variable’

⇒

‘No thing can be allocated, and not connected to a local variable’.

This is an invariant of the program as a whole.

Non-trivial global invariants can be preserved by defining assignment and destructors.